



ImageCaptioningModel structure (ViT encoder + attention decoder) — tensor shapes B=batch, T=caption length ≤ 32

Vision Encoder / ViT (patch=4, d_model=192, heads=3, ×12) cls=False



Text Embedding + DecoderBlock×1 + pred_head



Residual / skip edges — green dashed; image-feature feed — orange thick